

MURSHID · مرشد – THE COMPLETE GUIDE

Everything you need to explain the application: what it is, who it serves, how each screen works, what every element means, and how it connects to a real backend.

1 · WHAT MURSHID IS

Murshid (Arabic for "guide") is a multi-agent AI advisor for Emirati entrepreneurs, built for the **Khalifa Fund for Enterprise Development (KFED)**. Beneficiaries — small business owners like a home baker in Al Ain — talk to it over WhatsApp in the way they naturally speak: Emirati Arabic and English mixed mid-sentence, by text or voice note. Behind the conversation, four specialized AI agents cooperate to profile the entrepreneur, recommend funded KFED programs with real citations, watch for funding opportunities, and — crucially — hand the case to a human advisor the moment the AI is no longer confident it should answer alone.

The core pitch: *an AI that knows its limits*. Murshid doesn't pretend to be a lawyer or a banker. When a question crosses into territory like loan restructuring, it escalates to a human — and arrives with a bilingual briefing document already prepared, saving the advisor roughly two hours of prep.

2 · THE FOUR AGENTS

AGENT	WHAT IT DOES
Listener · المنصت steel blue	Intake. Transcribes bilingual voice notes (handling Arabic/English code-switching and Emirati dialect), classifies intent, and keeps the beneficiary's profile current — business type, location, stage, growth signals.
Planner · المخطط sand	Recommendation. Builds stepped program pathways from the KFED document corpus. Every recommendation carries a real citation — e.g. "Ruwad Program Guide, p.3" — so advisors can verify the source.
Scout · الكشف coral	Proactivity. Watches funding windows, deadlines and market signals, and fires <i>unprompted</i> alerts when a beneficiary matches — e.g. "Micro-financing window opens Monday — matches 4/5 criteria."
Safety Net · شبكة الأمان gold	Guardrails. Scores the system's confidence on every turn. Below the threshold (40%), it triggers the escalation: generates a bilingual Advisor Brief and hands the case to a human.

Each agent owns a consistent color everywhere in the interface, so an advisor can read the activity feed at a glance.

3 · SCREEN ONE — MISSION CONTROL

Murshid Mission Control.dc.html — the projected showcase screen (1920×1080, dark Gulf-night navy). It has three zones plus one theatrical moment:

Left — the beneficiary conversation

A WhatsApp-style chat with Fatima Al Mansoori: green-tinted bubbles for her messages, timestamps, read ticks, and a playable voice-note bubble with a waveform (press ► to "play" — the waveform fills). Messages mix Arabic (RTL) and English (LTR) in the same thread, rendered correctly via automatic direction handling. The Listener's transcript appears under the voice note.

Center — the agent activity feed

A live vertical timeline proving the system is real. Four agent chips at the top pulse while "thinking." Each event below shows the timestamp, which agent acted, what it did, and a one-line payload; Planner events carry citation chips (¶ Ruwad Program Guide, p.3). Events slide in as the story advances.

Right — the case panel

- **Profile card** — Fatima: home bakery, Al Ain, growth stage; fields fill in as the Listener works.
- **Recommended pathway** — a numbered 3-step track (Ruwad growth track → AED 50K micro-financing → F&B accelerator), each step cited.
- **Confidence meter** — the signature element: a thin gold arc with a large percentage. Steady gold while the system is sure.

The escalation takeover (the one theatrical moment)

When Fatima asks a complex loan-restructuring question, the arc visibly **drains from gold through amber to ember-red over ~3 seconds** (94% → 31%). Crossing the 40% threshold: the whole interface dims, the arc breaks open into a dashed ring, and a bilingual **Advisor Brief** slides up center-stage — English column left (LTR), Arabic column right (RTL): case summary, what's been tried, recommended next steps, and "prep time saved: ~2 hours."

Demo driver

→ / **Next beat** advances the 6-beat story · ← goes back · **A** runs the whole arc automatically in 60 seconds · **R** resets · **Esc** closes the brief. The story: standby → voice note arrives → Planner issues pathway → Scout alert → confidence drain → escalation.

4 · SCREEN TWO – THE WORKSPACE

Murshid Workspace.dc.html – the advisor's day-to-day back office (light, Industry design language). Seven sections, all navigable from the collapsible sidebar:

SECTION	CONTENTS
Overview	Usage vs. the 50,000-unit monthly limit, remaining allowance, task counts by status, a daily usage trend chart, most-used categories, recent agent activity, quick actions, and a warning banner when usage nears the limit.
Agents	One card per agent: role, description, live status, tasks completed, units consumed, error rate, capability tags, last action, and its three most recent tasks.
Tasks	Every task with search, status filter, category filter, sorting, and group-by-category. Clicking a row opens the full detail dialog: agent, beneficiary, timings, usage, output, files, activity timeline, errors – and actions (start, pause, continue, retry, cancel, duplicate, archive) with confirmations.
Categories	The five capability areas – Intake & Transcription, Beneficiary Profiling, Pathway Planning, Funding Scouting, Confidence & Escalation – each with its own usage vs. limit, owning agent, task counts and drill-down.
Activity / History	The complete audit trail: requests, agent actions, tool activity, outputs, files written, errors and retries, escalations – searchable, filterable by agent and entry type, paginated, with jump-to-task.
Usage & Limits	Deep metering: used/remaining/percent, billing period, projected exhaustion date, daily/weekly/monthly charts, usage by category, by agent, and by task (top consumers), plus a demo switcher showing the empty, loading, error, and limit-exceeded states.
Settings	Escalation threshold slider, usage-warning slider, notification toggles (Scout alerts, escalation emails, weekly digest), interface language (English/عربية), and history retention.

5 · DESIGN LANGUAGE

- **Mission Control palette:** deep Gulf-night navy ground; Khalifa Fund gold reserved almost exclusively for the confidence arc and escalation; warm sand for text/citations; one dawn-coral accent for Scout alerts. No decoration without meaning.

- **Typography:** IBM Plex Sans Arabic – a matched bilingual family – carries both scripts; a condensed geometric face for labels and the wordmark. Nothing under 14px; key numbers huge (readable from 6 meters).
- **Motion:** one orchestrated sequence (the takeover) plus restrained micro-motion (feed slide-ins, agent pulses). `prefers-reduced-motion` is respected throughout.
- **Structure is information:** agent colors are consistent everywhere; pathway steps are numbered because they truly are a sequence; citations look like citations.

6 • DATA & BACKEND INTEGRATION

Everything currently runs on **mock data – nothing is live**. The mocks are deliberately isolated so a real backend can replace them without touching the UI:

- **Mission Control:** a single `pollEvents()` function returns the messages, events, confidence and profile state for the current story beat. Swap its body for a FastAPI call with the same return shape.
- **Workspace:** `workspace-data.js` is the one mock module – agents, categories, tasks, activity and usage records plus a `fetchWorkspace()` function (with simulated latency for the loading state). Replace that one function with real endpoints.

7 • THE 60-SECOND PITCH

"Murshid is a WhatsApp-native AI advisor for Khalifa Fund entrepreneurs. Four agents work every conversation: the **Listener** understands mixed Arabic-English voice notes, the **Planner** recommends real KFED programs with page-level citations, the **Scout** proactively flags funding windows, and the **Safety Net** watches confidence on every answer. When a question is beyond the AI – watch the gold arc drain – it doesn't guess: it escalates to a human advisor with a bilingual brief already written, saving two hours of prep. The whole system is auditable: every task, every unit of usage, every escalation is tracked in the advisor workspace."